



# **Linux System Programming**

## **Part 3 - Filesystem and Files**

IBA Bulgaria  
2021



# Files in Linux

- “Everything is a file.”
- **inode** - data structure that describes a filesystem object.
- **Files** are always opened from user space by a **name**.
- A name and inode pair is called a **link**.
- **Regular files** - bytes of data.
- **Directories** - mapping between filenames and inodes (links).
- **Hard links** - multiple links map different names to the same inode.
- **Symbolic links** - like regular files which contain the complete pathname of the linked-to files.
- **Special files** - block device files, character device files, named pipes, Unix domain sockets.

# Filesystems and namespaces

- Linux provides a **global and unified namespace** of files and directories (with a root '/').
- A **filesystem** is a collection of files and directories in a formal and valid hierarchy.
- Filesystems may be individually added (**mounted**) to and removed (**unmounted**) from the global namespace of files and directories.
- Some special filesystems - '/dev' and '/proc'.

```
sysadmin@ubuntu:/dev$ ls
autofs          kmsg            sg0            tty30          tty61          ttyS5
block           kvm            sg1            tty31          tty62          ttyS6
bsg             lightnvm       shm            tty32          tty63          ttyS7
btrfs-control  log            snapshot      tty33          tty7           ttyS8
bus             loop0          snd            tty34          tty8           ttyS9
cdrom           loop1          sr0            tty35          tty9           ubuntu-vg
cdwr            loop2          stderr         tty36          ttyprintk     uhid
char            loop3          stdin          tty37          ttyS0          uinput
console         loop4          stdout         tty38          ttyS1          urandom
core            loop5          tty            tty39          ttyS10         userio
cpu             loop6          tty0           tty4            ttyS11         vcs
cpu_dma_latency loop7           tty1           tty40           ttyS12         vcs1
cuse            loop-control   tty10          tty41           ttyS13         vcs2
disk            mapper         tty11          tty42           ttyS14         vcs3
dm-0            mcelog         tty12          tty43           ttyS15         vcs4
```

```
sysadmin@ubuntu:/proc$ ls
1      18      25210  42     59      766      diskstats  net
10     18956   25211  43     6        774      dma        pagetypeinfo
1044   19      25212  431    60       78       driver     partitions
1047   2       25298  432    61       783      execdomains sched_debug
11     20      25299  437    62       7922     fb         schedstat
12     20693  25327  439    63       8        filesystems scsi
124    21      255    44     64       80       fs         self
125    21753  256    440    65       800      interrupts slabinfo
128    21816  26     441    66       81       iomem      softirqs
12841  21934  27     442    67       815      ioports    stat
129    22     27640  445    68       82       ipmi       swaps
13     23259  28     447    69       822      irq        sys
130    24     29998  448    7        855      kallsyms   sysrq-trigger
131    240    30     449    70       882      kcore      sysvipc
```



# Working with Files in C

- Before a file can be read from or written to, it must be **opened**.
- Each open instance of a file is given a unique **file descriptor (fd)**.
- File descriptors are represented by the C *int* type.
- A single file can be opened more than once, by a different or even the same process.
- **open()** / **fopen()**- opens file and returns its file descriptor.
- **read()** / **fread()** - reads data from a file.
- **write()** / **fwrite()** - writes data into a file.
- **fflush()** - flushes a stream.
- **close()** - unmaps the a file descriptor with the associated file.
- **lseek()** / **fseek()** - set the file position of a file descriptor to a given value.
- **fcntl()** - manipulate file descriptor, for example for locking.
- **errno** - number of last error.



# Buffered vs unbuffered streams


- **Stream** is a representation of flow of data from one side to another e.g. from disk to memory and from memory to disk.
- **File** is a representation to store data on disk file. File uses streams to store and load data.
- **Buffer** is (often) used to hold stream data temporarily.
- Characters written to or read from an **unbuffered** stream are transmitted individually to or from the file as soon as possible.
- Characters written to or read from a **fully buffered** stream are transmitted to or from the file in blocks of arbitrary size.
- Characters written to a **line buffered** stream are transmitted to the file in blocks when a newline character is encountered.

# Buffered vs unbuffered streams

Low-level file routines (**unbuffered** streams):

- `open()`, `read()`, `write()`, `lseek()`, etc.
- Part of `<unistd.h>` library.
- Work with file descriptors of type `int`.
- Treat the input/output as **binary** data.


## Unbuffered:



High-level file routines (**buffered** streams):

- `fopen()`, `fread()`, `fwrite()`, `fseek()`, etc.
- Part of `<stdio.h>` library.
- Use the object type `FILE`.
- Treat the input/output as **text** streams.
- The read/write of the accumulated buffer can be forced with `fflush()`.

## Buffered:



# Basic file processing

```
#include <unistd.h>
int open (const char *name, int flags);
int open (const char *name, int flags, mode_t mode);

int close (int fd);
```

If **flags** has **O\_CREAT** set, the **mode** is a bit mask of:

S_IRWXU	S_IWGRP
S_IRUSR	S_IXGRP
S_IWUSR	S_IRWXO
S_IXUSR	S_IROTH
S_IRWXG	S_IWOTH
S_IRGRP	S_IXOTH

**flags** parameter is a bit mask of the following bits:

O_APPEND	The file will be opened in append mode.
O_ASYNC	<b>SIGIO</b> generated when readable or writable.
O_CREAT	If the file doesn't exist - create it.
O_DIRECT	Opened for direct I/O.
O_DIRECTORY	If <b>name</b> is not a directory, <b>open( )</b> will fail.
O_EXCL	If O_CREAT and file exists, <b>open( )</b> will fail.
O_LARGEFILE	A file larger than 2G to be opened.
O_NOCTTY	This flag is not frequently used.
O_NOFOLLOW	If <b>name</b> is a symbolic link, <b>open( )</b> will fail.
O_NONBLOCK	If possible, open in nonblocking mode.
O_SYNC	The file will be opened for synchronous I/O.
O_TRUNC	If the file exists, truncated it to zero length.

# Open and close a file

Initialize the variables

Open the file as read-only and get the **fd**

If the **fd** is **-1**, print error message

Otherwise, print the value of **fd** (success)

Close the file

If the close failed, print error message

```
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>

int main(int argc, char * argv[]) {
    int fd;
    fd = open ("/proc/self/environ", O_RDONLY);
    if (fd == -1) {
        printf("ERROR opening 'environ'!\n");
    } else {
        printf("File Descriptor of 'environ' = %d\n", fd);
    }
    // write/read commands to be added here
    if (close (fd) == -1)
        printf("ERROR closing the file!\n");
}
```





## Reading file contents

```
#include <unistd.h>
ssize_t read (int fd, void *buf, size_t len);
```

- Each call reads up to **len** bytes into **buf** from the current file offset of the file referenced by **fd**.
- On success, the number of bytes written into **buf** is returned.
- On error, the call returns **-1** and **errno** is set.

# Read and print a file

readfile.c

```
#define BUF_SIZE 1000000
...
int fd;
ssize_t len;
char buf[BUF_SIZE];
```

Initialize the variables

Open './readfile.c', if failed print error & exit

While reading from file returns length  $\neq 0$

If length is -1 and errno is EINTR, try to read again

If length is -1 and errno  $\neq$  EINTR, print error and finish reading

Otherwise print the buffer

Close the file and notify on error

```
while ((len = read(fd, buf, BUF_SIZE - 1)) != 0)
{
    if (len == -1)
    {
        if (errno == EINTR) continue;
        printf("ERROR reading the file!\n");
        break;
    }
    printf("%s", buf);
}
```



## Create and write into a file

```
#include <unistd.h>
```

```
ssize_t write (int fd, const void *buf, size_t count);
```

- Writes up to **count** bytes starting at **buf** to the current file position of the file referenced by the file descriptor **fd**.
- On success, the number of bytes written is returned, and the file position is updated.
- On error, the call returns **-1** and **errno** is set.

writefile.c

# Write sentences into file

Initialize the variables

Create empty './sentences.txt', if failed  
print error & exit

Do 100 times

Call **getSentence()** to get a new  
text into the buffer and its length

Write the buffer into the file and get  
the number of bytes written

If the number of written bytes is -1,  
print error and stop writing.

Close the file and notify the error or success

```
fd = open("./sentences.txt", O_WRONLY | O_CREAT |  
O_TRUNC, 0644);
```

```
int getSentence(char *buf)  
...  
for (int i=0; i < 100; i++)  
{  
    int text_len = getSentence(buf);  
    ssize_t nr = write(fd, buf, text_len);  
    if (nr == -1)  
    {  
        printf("ERROR writing to the file!\n");  
        break;  
    }  
}
```



## Seeking in files and Sparse files

```
#include <stdio.h>
FILE *fopen(const char *pathname, const char *mode);
size_t fread(void *ptr, size_t size, size_t nmem, FILE *stream);
int fseek(FILE *stream, long offset, int whence);
int fclose(FILE *stream);
```

- The behavior of **fseek()** depends on the *origin* argument: **SEEK\_CUR**, **SEEK\_END** or **SEEK\_SET**.
- The call returns the new file position on success.
- On error, the call returns **-1** and **errno** is set.
- Seeking after the end of a file and then writing in it causes **holes** padded with zeros.
- Files with holes are called **sparse files**.
- Holes do not occupy any physical disk space.
- **du** - estimate file space usage.

# Make a file with a hole



Initialize the variables

Create a file, its name is the first argument passed to the program

If failed to open - notify error and exit

Write the name of the file into the file

Jump 16777216 bytes forward into the file

Write the name of the file into the file again

Close the file

```
#include <stdio.h>
#include <string.h>

#define BIG_SIZE 0x1000000

int main(int argc, char * argv[])
{
    FILE * f;
    f = fopen(argv[1], "w");
    if (f == NULL)
    {
        printf("ERROR creating file: %s", argv[1]);
        return 1;
    }
    fwrite(argv[1], 1, strlen(argv[1]), f);
    fseek(f, BIG_SIZE, SEEK_CUR);
    fwrite(argv[1], 1, strlen(argv[1]), f);
    fclose(f);
}
```

# Locking files

```
#include <sys/types.h>
#include <fcntl.h>
#include <unistd.h>

off_t lseek(int fd, off_t offset, int whence);
int fcntl(int fd, int cmd, ... /* arg */ );
```

```
struct flock {
    ...
    short l_type;    /* Type of lock: F_RDLCK,
                     F_WRLCK, F_UNLCK */
    short l_whence;  /* How to interpret l_start:
                     SEEK_SET, SEEK_CUR, SEEK_END */
    off_t l_start;   /* Starting offset for lock */
    off_t l_len;     /* Number of bytes to lock */
    pid_t l_pid;     /* PID of process blocking our lock
                     (set by F_GETLK and F_OFD_GETLK) */
    ...
};
```

Use the **fcntl()** call providing a pointer to a **flock** structure. This call manipulates the file descriptor **fd**, depending on the command **cmd**:

- To lock a block of a file use **F\_SETLK**.
- If the block is already locked use **F\_GETLK** to get information about the locking process.
- To unlock a block of a file use **F\_SETLK** but set **flock.l\_type = F\_UNLCK**.

# Lock and write in there

Initialize the variables

Open for writing or create './testlocks.txt'

While lock of 64 bytes from **offset** fails:

Get information about the locking process and print info about it

Move the **offset** 64 bytes further

Move **offset** bytes from the start of the file and write info about the current process

Wait for **Enter** key press

Unlock the 64 locked bytes, notify if failed

Close the file

```
struct flock fi;
fi.l_type = F_WRLCK;
fi.l_whence = SEEK_SET;
fi.l_start = 0;
fi.l_len = 64;
```

```
fd = open("testlocks.txt", O_RDWR|O_CREAT);
while (fcntl(fd, F_SETLK, &fi) == -1)
{
    fcntl(fd, F_GETLK, &fi);
    printf("bytes %i - %i locked by process %i\n", off, off+64,
        fi.l_pid);
    off += 64;
    fi.l_start = off;
}
lseek(fd, off, SEEK_SET);
sprintf(str, "Stored by process %i", getpid());
write(fd, str, strlen(str));
getchar();
fi.l_type = F_UNLCK;
if (fcntl(fd, F_SETLK, &fi) == -1)
    printf("ERROR while locking!\n");
close(fd);
```



# Exercise



## Program *FileManipulator*:

Write a program ('**fmanipulator.c**'), which takes 3 arguments: **words\_count**, **min\_length**, **max\_length**. The program should generate a file called '**words.txt**', which contains **words\_count** words (separated with spaces). To generate the words use the following function from the code files in '**/day03/rndword/**' (note that the function will not generate real words, but rather random sequences of characters):

```
char * rndword(int min_length, int max_length);
```

Then the program should print the generated file to the screen.

After the program creates '**words.txt**' and before writing into it it should lock the first 100 bytes of the file. When the program prints out the generated words it should wait for **Enter** key press and then unlock the file. Respectively, if the file was already locked by another process, the program should notify about this and wait for **Enter** key press, before trying again (to lock, generate and write).

Compile and run a couple of instances to test.